

RLearning: Short guides to reinforcement learning

Unit 3-1: Overview Reinforcement Learning

Davud Rostam-Afschar (Uni Mannheim)

How computers (humans) learn?


Markov Decision Process

► Definition

- States: $s \in S$
- Actions: $a \in A$
- Rewards: $r \in \mathbb{R}$
- Transition model: $\mathbb{P}(s_t | s_{t-1}, a_{t-1})$
- Reward model: $\mathbb{P}(r_t | s_t, a_t)$
- Discount factor: $0 \leq \gamma \leq 1$
 - discounted: $\gamma < 1$
 - undiscounted: $\gamma = 1$
- Horizon (i.e., # of time steps): h
 - Finite horizon: $h \in \mathbb{N}$
 - infinite horizon: $h = \infty$
- Goal: find optimal policy π^* such that

$$\pi^* = \operatorname{argmax}_{\pi} \sum_{t=0}^h \gamma^t \mathbb{E}_{\pi} [r_t]$$

Reinforcement Learning Problem



Goal: Learn to choose actions that maximize rewards

Reinforcement Learning

- ▶ Definition
 - ▶ States: $s \in S$
 - ▶ Actions: $a \in A$
 - ▶ Rewards: $r \in \mathbb{R}$
 - ▶ Transition model: $\mathbb{P}(s_t | s_{t-1}, a_{t-1})$ unknown
 - ▶ Reward model: $\mathbb{P}(r_t | s_t, a_t)$ model
 - ▶ Discount factor: $0 \leq \gamma \leq 1$
 - ▶ discounted: $\gamma < 1$
 - ▶ undiscounted: $\gamma = 1$
 - ▶ Horizon (i.e., # of time steps): h
 - ▶ Finite horizon: $h \in \mathbb{N}$
 - ▶ infinite horizon: $h = \infty$
- ▶ Goal: find optimal policy π^* such that

$$\pi^* = \operatorname{argmax}_{\pi} \sum_{t=0}^h \gamma^t \mathbb{E}_{\pi} [r_t]$$

Policy Optimization

- ▶ **Markov Decision Process:**

- ▶ Find optimal policy given transition and reward model
- ▶ Execute policy found

- ▶ **Reinforcement learning:**

- ▶ Learn an optimal policy while interacting with the environment


Policy Optimization

► **Markov Decision Process:**

- ▶ Find optimal policy given transition and reward model
- ▶ Execute policy found

► **Reinforcement learning:**

- ▶ Learn an optimal policy while interacting with the environment





Policy Optimization

► Markov Decision Process:

- ▶ Find optimal policy given transition and reward model
- ▶ Execute policy found


► Reinforcement learning:

- ▶ Learn an optimal policy while interacting with the environment



Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE


- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

Example: MENACE

- ▶ Machine Educable Noughts And Crosses Engine (MENACE)
- ▶ Build by Donald Michie in 1961 using 304 matchboxes



Source: Matthew Scroggs

- | | | |
|-------------|---|-------------------------------|
| Lose | → | Remove the bead from each box |
| Win | → | Add three beads to each box |
| Draw | → | Add one bead to each box |

Important Components in Reinforcement Learning

Reinforcement learning agents may or may not include the following components:

- ▶ **Model:** $\mathbb{P}(s' | s, a), \mathbb{P}(r | s, a)$
 - ▶ Environment dynamics and rewards
- ▶ **Policy:** $\pi(s)$
 - ▶ Agent action choices
- ▶ **Value function:** $V(s)$
 - ▶ Expected total rewards of the agent's policy

Important Components in Reinforcement Learning

Reinforcement learning agents may or may not include the following components:

- ▶ **Model:** $\mathbb{P}(s' | s, a), \mathbb{P}(r | s, a)$
 - ▶ Environment dynamics and rewards
- ▶ **Policy:** $\pi(s)$
 - ▶ Agent action choices
- ▶ **Value function:** $V(s)$
 - ▶ Expected total rewards of the agent's policy
- ▶ **Quality function:** $Q(s, a)$
 - ▶ Expected total rewards of taking a specific action in a given state

Bellman's Equation

- ▶ Optimal state value function $V^*(s)$

$$V^*(s) = \max_a E[r | s, a] + \gamma \sum_{s'} \mathbb{P}(s' | s, a) V^*(s')$$

- ▶ Optimal state-action value function $Q^*(s, a)$

$$Q^*(s, a) = E[r | s, a] + \gamma \sum_{s'} \mathbb{P}(s' | s, a) \max_{a'} Q^*(s', a')$$

where $V^*(s) = \max_a Q^*(s, a)$

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

Categorizing RL Agents

Categorizing RL Agents

Value based

- ▶ No policy (implicit)
- ▶ Value function

Policy based

- ▶ Policy
- ▶ No value function

Actor critic

- ▶ Policy
- ▶ Value function

Model based

- ▶ Transition and reward model

Model free

- ▶ No transition model (implicit)
- ▶ No reward model (implicit)

Categorizing RL Agents

Value based

- ▶ No policy (implicit)
- ▶ Value function

Policy based

- ▶ Policy
- ▶ No value function

Actor critic

- ▶ Policy
- ▶ Value function

Model based

- ▶ Transition and reward model

Model free


- ▶ No transition model (implicit)
- ▶ No reward model (implicit)

Imitation Learning

- ▶ No transition model (implicit)
- ▶ Reward model implicit through experts

Dynamic Programming Backup

$$V(S_t) \leftarrow \mathbb{E}_\pi [R_{t+1} + \gamma V(S_{t+1})]$$




Source: David Silver

RL Algorithms

Monte Carlo Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$




Source: David Silver

RL Algorithms

Temporal Difference Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



Source: David Silver

Toy Maze Example

3	r	r	r	
2	u		u	
1	u			
	1	2	3	4

Start state: (1,1)

Terminal states: (4,2), (4,3)

No discount: $\gamma=1$

Reward is -0.04 for non-terminal states

Four actions:

- ▶ up (**u**),
- ▶ left (**|**),
- ▶ right (**r**),
- ▶ down (**d**)

Do not know the transition probabilities

What is the value $V(s)$ of being in state s

Toy Maze Example (No Learning, Noise 20%)



References |

- DENERO, J., D. KLEIN, B. MILLER, N. HAY, AND P. ABBEEL (2013): "The Pacman AI Projects," <http://inst.eecs.berkeley.edu/~cs188/pacman/>, Developed at UC Berkeley. Core by John DeNero and Dan Klein; student autograding by Brad Miller, Nick Hay, and Pieter Abbeel.
- POUPART, P. (2025): "Pascal Poupart's Homepage,"
<https://cs.uwaterloo.ca/~ppoupart/>, Accessed: 2025-05-24.
- RUSSELL, S. J., AND P. NORVIG (2016): *Artificial intelligence: a modern approach*. Pearson.
- SIGAUD, O., AND O. BUFFET (2013): *Markov decision processes in artificial intelligence*. John Wiley & Sons.
- SUTTON, R. S., AND A. G. BARTO (2018): "Reinforcement learning: An introduction," *A Bradford Book*, Available at
<http://incompleteideas.net/book/the-book-2nd.html>.
- SZEPESVÁRI, C. (2022): *Algorithms for reinforcement learning*. Springer nature, Available at <https://sites.ualberta.ca/~szepesva/RLBook.html>.

Takeaways

Takeaways

- ▶ "By three methods we may learn wisdom:
 - ▶ First, by reflection, which is ~~noblest~~ model-based RL;
 - ▶ Second, by imitation, which is easiest imitation learning; and
 - ▶ third by experience, which is the ~~bitterest~~ model free RL."
- ▶ RL agent types:
 - ▶ value-based,
 - ▶ policy-based,
 - ▶ value-policy-based (actor-critic),
 - ▶ model-based,
 - ▶ model-free
 - ▶ imitation learning